

Data preparation: exploration and normalization

- The start point was to observe and analyze the tables in the dataset to understand how the data is distributed (Pandas libraries for Python is used to handle the tables as dataframes).
- The common functions of Pandas DataFrame to use for an overview of the table are the following:

```
dataframe.info()  
dataframe.head()  
dataframe.describe()
```

- After reviewing the contents of the various tables in the MIMIC database, only some tables were selected and loaded into DataFrames using Pandas because it was assumed that it wasn't necessary or useful for prediction to include all tables.

So the table selected that compose the baseline dataset are the following:

- **ADMISSIONS.csv** that defines a patient's hospital admission,
- **PATIENTS.csv** that defines a single patient,
- **DIAGNOSES_ICD.csv** that contains ICD diagnoses for patients, most notably ICD-9 diagnoses (after this matter will be better explained),

Other tables will be considered, if necessary for the specific task.

```
In [1]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt
```

ADMISSIONS table

```
In [2]: mimic4_path = './data/mimic-iv-clinical-database-demo-2.2/'  
  
# read admissions table  
def read_admissions_table(mimic4_path):  
    admits = pd.read_csv(mimic4_path + 'hosp/admissions.csv')  
    # Pre-emptively don't include some columns that I don't need  
    admits = admits[['subject_id', 'hadm_id', 'admittime', 'dischtime', 'deathtime',  
    # Converts dates to a proper format  
    admits.admittime = pd.to_datetime(admits.admittime)  
    admits.dischtime = pd.to_datetime(admits.dischtime)  
    admits.deathtime = pd.to_datetime(admits.deathtime)  
    return admits
```

```
admits = read_admissions_table(mimic4_path)
admits.head()
```

Out[2]:

	subject_id	hadm_id	admittime	dischtime	deathtime	admission_type	insurance
0	10004235	24181354	2196-02-24 14:38:00	2196-03- 04 14:02:00	NaT	URGENT	Medicaid
1	10009628	25926192	2153-09-17 17:08:00	2153-09- 25 13:20:00	NaT	URGENT	Medicaid
2	10018081	23983182	2134-08-18 02:02:00	2134-08- 23 19:35:00	NaT	URGENT	Medicare
3	10006053	22942076	2111-11-13 23:39:00	2111-11- 15 17:20:00	2111-11-15 17:20:00	URGENT	Medicaid
4	10031404	21606243	2113-08-04 18:46:00	2113-08- 06 20:57:00	NaT	URGENT	Other

In [3]: `admits.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 275 entries, 0 to 274
Data columns (total 7 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   subject_id      275 non-null    int64  
 1   hadm_id         275 non-null    int64  
 2   admittime       275 non-null    datetime64[ns]
 3   dischtime       275 non-null    datetime64[ns]
 4   deathtime       15 non-null     datetime64[ns]
 5   admission_type 275 non-null    object  
 6   insurance        275 non-null    object  
dtypes: datetime64[ns](3), int64(2), object(2)
memory usage: 15.2+ KB
```

When DEATHTIME in ADMISSIONS is not null then the patient associated died at the hospital, so we mark this distinction with a boolean variable.

In [4]: `admits['died_at_the_hospital'] = admits['deathtime'].notnull().map({True:1, False:0})`

In [5]: `# ADMISSION_TYPE`
`admits['admission_type'].value_counts()`

```
Out[5]: admission_type
EW EMER.          104
OBSERVATION ADMIT    45
URGENT            38
EU OBSERVATION     30
SURGICAL SAME DAY ADMISSION 18
DIRECT EMER.       15
ELECTIVE          13
DIRECT OBSERVATION   7
AMBULATORY OBSERVATION 5
Name: count, dtype: int64
```

Comprese into EMERGENCY

The category URGENT is a lot similar semantically to EMERGENCY or DIRECT EMERGENCY, so could combine these categories in EMERGENCY. We could do the same process also for all those categories that are related to OBSERVATION.

```
In [6]: admits.loc[:, 'admission_type'] = admits['admission_type'].replace('EW EMER.', 'EMERGENCY')
admits.loc[:, 'admission_type'] = admits['admission_type'].replace('DIRECT EMER.', 'EMERGENCY')
admits.loc[:, 'admission_type'] = admits['admission_type'].replace('URGENT', 'EMERGENCY')

admits['admission_type'].value_counts()
```

```
Out[6]: admission_type
EMERGENCY          157
OBSERVATION ADMIT    45
EU OBSERVATION     30
SURGICAL SAME DAY ADMISSION 18
ELECTIVE          13
DIRECT OBSERVATION   7
AMBULATORY OBSERVATION 5
Name: count, dtype: int64
```

```
In [7]: # read patients table
def read_patients_table(mimic4_path):
    pats = pd.read_csv(mimic4_path + 'hosp/patients.csv')
    # Pre-emptively don't include some columns that I don't need
    pats = pats[['subject_id', 'gender', 'anchor_age', 'dod']]
    pats.dod = pd.to_datetime(pats.dod)
    return pats

patients = read_patients_table(mimic4_path)
patients.head()
```

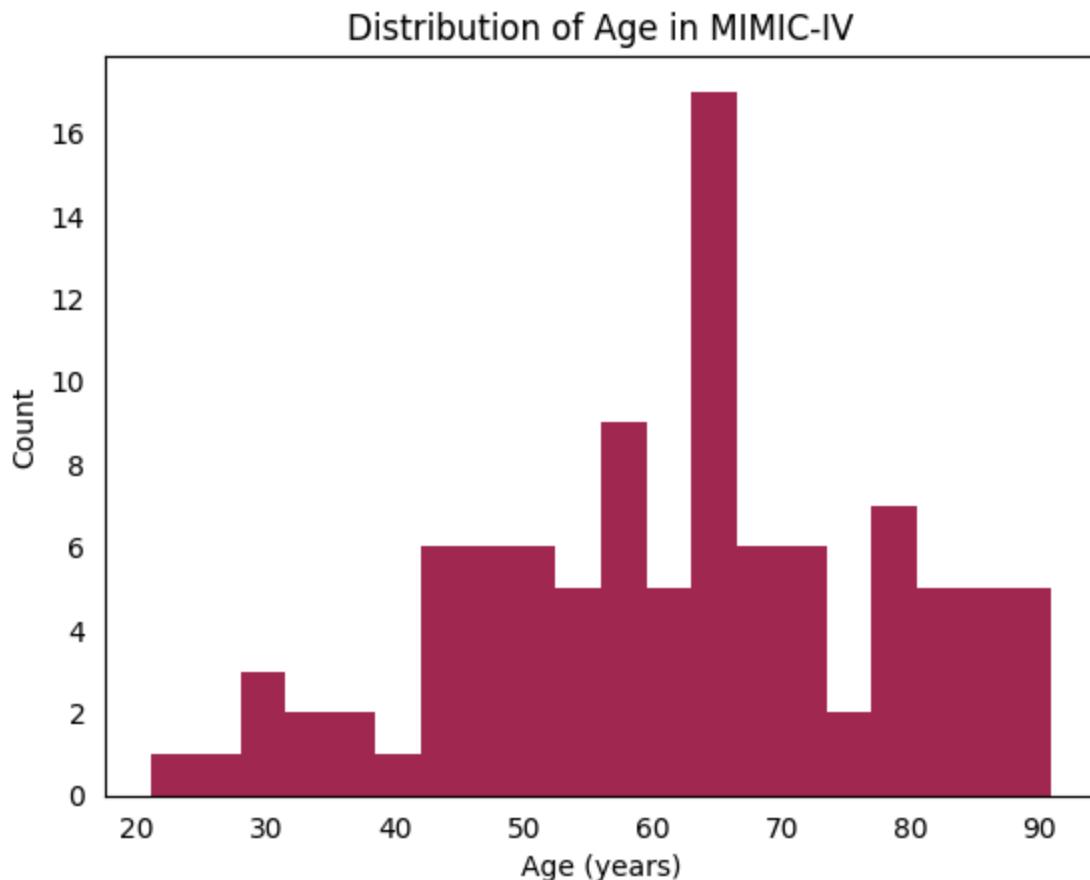
Out[7]:

	subject_id	gender	anchor_age	dod
0	10014729	F	21	NaT
1	10003400	F	72	2137-09-02
2	10002428	F	80	NaT
3	10032725	F	38	2143-03-30
4	10027445	F	48	2146-02-09

As we can see from age distribution, patients in their childhood are not present, this reflects the fact that MIMIC-IV as MIMIC-III does not contain data from pediatric patients.

In [8]:

```
#Let's see the distribution of age
plt.hist(patients['anchor_age'], bins=20, color="#a12a52")
plt.ylabel('Count')
plt.xlabel('Age (years)')
plt.title('Distribution of Age in MIMIC-IV')
plt.tick_params(left=False, bottom=False, top=False, right=False)
plt.show()
```

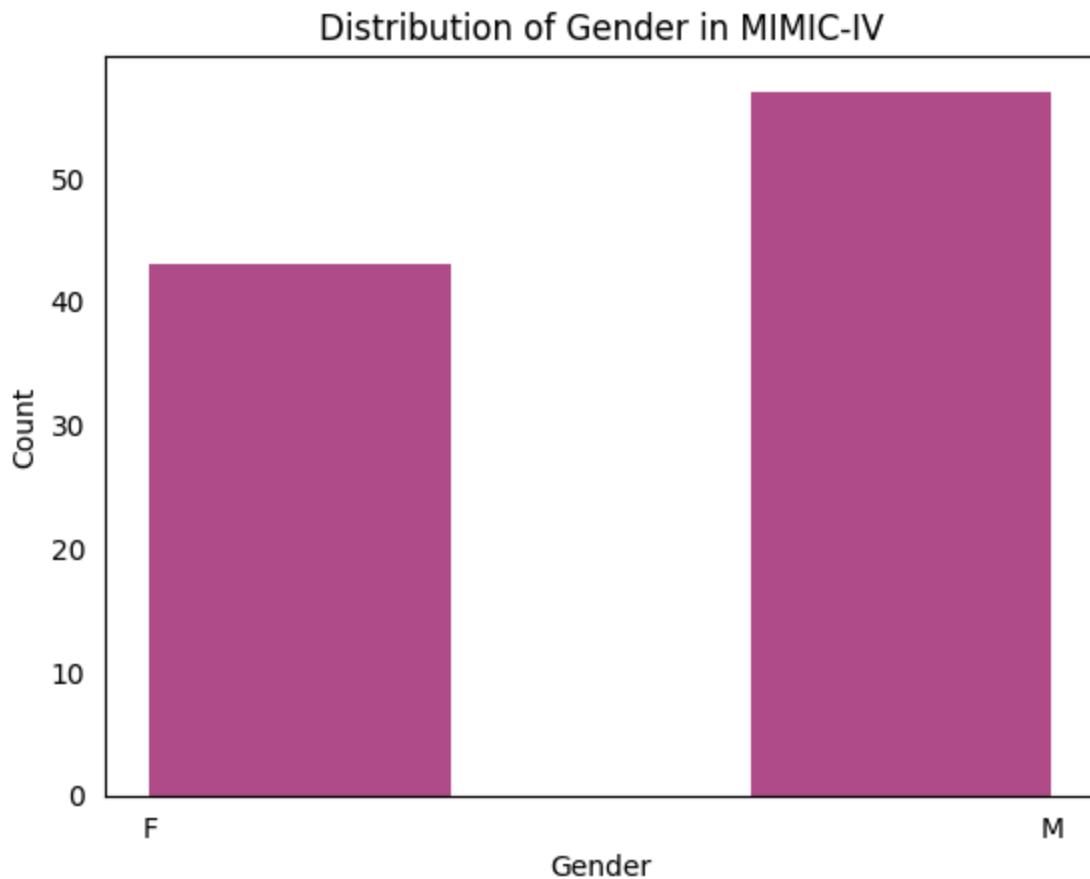


In [9]:

```
patients['gender'].value_counts()
```

```
Out[9]: gender  
M      57  
F      43  
Name: count, dtype: int64
```

```
In [10]: # distribution of gender  
plt.hist(patients['gender'], bins=3, color='#b24b88')  
plt.ylabel('Count')  
plt.xlabel('Gender')  
plt.title('Distribution of Gender in MIMIC-IV')  
plt.tick_params(left=False, bottom=False, top=False, right=False)  
plt.show()
```



We can see how the gender is almost equally balanced.

Now we merge patients and admissions tables on 'subject_id' link.

```
In [11]: # merge the PATIENTS table with ADMISSIONS table  
admits_patients = pd.merge(admits, patients, how='inner', on='subject_id')  
admits_patients.head()
```

Out[11]:

	subject_id	hadm_id	admittime	dischtime	deathtime	admission_type	insurance	died
0	10004235	24181354		2196-02-24 14:38:00	2196-03-04 14:02:00	NaT	EMERGENCY	Medicaid
1	10009628	25926192		2153-09-17 17:08:00	2153-09-25 13:20:00	NaT	EMERGENCY	Medicaid
2	10018081	23983182		2134-08-18 02:02:00	2134-08-23 19:35:00	NaT	EMERGENCY	Medicare
3	10006053	22942076		2111-11-13 23:39:00	2111-11-15 17:20:00	2111-11-15 17:20:00	EMERGENCY	Medicaid
4	10031404	21606243		2113-08-04 18:46:00	2113-08-06 20:57:00	NaT	EMERGENCY	Other

read diagnoses_icd table

In [12]:

```
def read_diagnoses_icd_table(mimic4_path):
    diag_icds = pd.read_csv(mimic4_path + 'hosp/diagnoses_icd.csv')
    return diag_icds

diag_icds = read_diagnoses_icd_table(mimic4_path)
diag_icds.head()
```

Out[12]:

	subject_id	hadm_id	seq_num	icd_code	icd_version
0	10035185	22580999	3	4139	9
1	10035185	22580999	10	V707	9
2	10035185	22580999	1	41401	9
3	10035185	22580999	9	3899	9
4	10035185	22580999	11	V8532	9

In [13]:

```
diag_icds.count()
```

Out[13]:

subject_id	4506
hadm_id	4506
seq_num	4506
icd_code	4506
icd_version	4506
	dtype: int64

In [14]:

```
diag_icds.icd_version.value_counts()
```

```
Out[14]: icd_version
10    2313
9     2193
Name: count, dtype: int64
```

We can notice that there are two version of icd_code: version 9 and 10. In general, ICD-10 codes are more detailed, but they could be mapped and converted to ICD-9 because they express the same concept.

Since we are dealing with a dataframe with a lot of entries and in any case it should later be reduced in size for a faster test, for simplicity we consider only the diagnoses with ICD9 codes.

```
In [15]: diag_icds = diag_icds[diag_icds['icd_version'] == 9]
diag_icds.icd_version.value_counts()
```

```
Out[15]: icd_version
9     2193
Name: count, dtype: int64
```

International Classification of Diseases, Clinical Modification (ICD-CM in version 9 and 10) is an adaption created by the U.S. National Center for Health Statistics (NCHS) and used in assigning diagnostic and procedure codes associated with inpatient, outpatient, and physician office utilization in the United States.

```
In [16]: print('There are {} unique ICD9 codes in this dataset.'.format(diag_icds['icd_code']))
There are 740 unique ICD9 codes in this dataset.
```

```
In [17]: # Filter out E and V codes from ICD9 codes since processing will be done on the num
import warnings
warnings.filterwarnings('ignore') # To ignore all warnings

diag_icds['recode'] = diag_icds['icd_code']
diag_icds.loc[:, 'recode'] = diag_icds['recode'].fillna(value='999')

diag_icds['recode'] = diag_icds['recode'][~diag_icds['recode'].str.contains("[a-zA-Z]")]
diag_icds['recode'].fillna(value='999', inplace=True)

# Take in consideration just the first 3 integers of the ICD9 code
diag_icds['recode'] = diag_icds['recode'].str.slice(start=0, stop=3, step=1)
diag_icds['recode'] = diag_icds['recode'].astype(int)
diag_icds.head()
```

Out[17]:

	subject_id	hadm_id	seq_num	icd_code	icd_version	recode
0	10035185	22580999	3	4139	9	413
1	10035185	22580999	10	V707	9	999
2	10035185	22580999	1	41401	9	414
3	10035185	22580999	9	3899	9	389
4	10035185	22580999	11	V8532	9	999

In [18]:

```
# ICD-9 Main Category ranges
icd9_ranges = [(1, 140), (140, 240), (240, 280), (280, 290), (290, 320), (320, 390),
                (390, 460), (460, 520), (520, 580), (580, 630), (630, 680), (680, 71
                (710, 740), (740, 760), (760, 780), (780, 800), (800, 1000), (1000,
                (1000, 1040)

# Associated category names
diag_dict = {0: 'infectious', 1: 'neoplasms', 2: 'endocrine', 3: 'blood',
             4: 'mental', 5: 'nervous', 6: 'circulatory', 7: 'respiratory',
             8: 'digestive', 9: 'genitourinary', 10: 'pregnancy', 11: 'skin',
             12: 'muscular', 13: 'congenital', 14: 'prenatal', 15: 'misc',
             16: 'injury', 17: 'misc'}

# Re-code in terms of integer
for num, cat_range in enumerate(icd9_ranges):
    diag_icds['recode'] = np.where(diag_icds['recode'].between(cat_range[0],cat_ran
        (cat_range[1]),True))

# Convert integer to category name using diag_dict
diag_icds['super_category'] = diag_icds['recode'].replace(diag_dict)
diag_icds.head()
```

Out[18]:

	subject_id	hadm_id	seq_num	icd_code	icd_version	recode	super_category
0	10035185	22580999	3	4139	9	6	circulatory
1	10035185	22580999	10	V707	9	16	injury
2	10035185	22580999	1	41401	9	6	circulatory
3	10035185	22580999	9	3899	9	5	nervous
4	10035185	22580999	11	V8532	9	16	injury

In [19]:

```
# Create list of diagnoses for each admission
hadm_list = diag_icds.groupby('hadm_id')['super_category'].apply(list).reset_index()
hadm_list.head()
```

Out[19]:

	hadm_id	super_category
0	20044587	[injury, injury, circulatory, circulatory, end...
1	20192635	[endocrine, digestive, injury, misc, genitouri...
2	20214994	[injury, circulatory, injury, mental, injury, ...
3	20285402	[infectious, infectious, injury, infectious, e...
4	20291550	[injury, injury, genitourinary, respiratory, e...

In [20]:

```
stacked = hadm_list['super_category'].apply(pd.Series).stack()
hadm_item = pd.get_dummies(stacked).groupby(level=0).sum()
hadm_item.head()
```

Out[20]:

	blood	circulatory	congenital	digestive	endocrine	genitourinary	infectious	injury	mental
0	0	3	0	0	1	0	0	0	3
1	0	1	1	1	3	3	0	0	4
2	2	4	0	3	2	2	1	1	17
3	1	2	0	1	3	0	3	1	2
4	0	0	0	0	1	1	1	1	6

Join back with HADM_ID

In [21]:

```
hadm_item = hadm_item.join(hadm_list['hadm_id'], how="outer")
```

Merge with main dataframe

In [22]:

```
admits_patients_diag = pd.merge(admits_patients, hadm_item, how='inner', on='hadm_id')
admits_patients_diag.head()
```

Out[22]:

	subject_id	hadm_id	admittime	dischtime	deathtime	admission_type	insurance	died
0	10004235	24181354	2196-02-24 14:38:00	2196-03-04 14:02:00	NaT	EMERGENCY	Medicaid	
1	10009628	25926192	2153-09-17 17:08:00	2153-09-25 13:20:00	NaT	EMERGENCY	Medicaid	
2	10018081	23983182	2134-08-18 02:02:00	2134-08-23 19:35:00	NaT	EMERGENCY	Medicare	
3	10006053	22942076	2111-11-13 23:39:00	2111-11-15 17:20:00	2111-11-15 17:20:00	EMERGENCY	Medicaid	
4	10005817	20626031	2132-12-12 01:43:00	2132-12-20 15:04:00	NaT	EMERGENCY	Medicare	

5 rows × 27 columns

**Save this version of the dataframe to a csv.
It will be used as baseline for our
prediction tasks**

In [23]: `admits_patients_diag.to_csv('admits_patients_diag.csv')`